

Exponential Smoothing to Analyze End of Summer Trends

Alena Fedash

2022-09-19

Contents

Summary of Results	1
Solution in R	2

Summary of Results

Since exponential smoothing is a time series forecasting method, I began my analysis by converting the data to a **time series object**. As frequency, I used the number of days which we had for each year in the data set (July through October), and the start value was the first year we had, 1996. This way, I generated a times series object with 20 ‘seasons’, each season being defined by temperatures from July 1 to October 31.

After that, I performed visual inspection of the times series, and noticed that, according to the plot, the data is more likely to have an **additive seasonality factor**, rather than multiplicative, since there was no obvious funnel-like shape on the graph.

Nevertheless, I decomposed the time series data trying both methods, multiplicative and additive seasonality. Both showed similar results, the only difference being that in multiplicative version random and seasonal components were displayed as coefficient rather than additive values. **Decomposition also showed that the trend factor is rather insignificant**, suggesting there is no year-to-year trend in temperature change. This is an important signal that there is no general ‘warming’ in our data that is becoming more obvious every year - the temperatures are rather stable. There was also a **significant random component**, which had a wider value range compared to seasonality (and trend). That is a sign that it would be hard to build accurate predictions (if needed) with such data, as there is too much ‘noise’ present. This is logical, since temperature prognosis is usually done using winds, cyclones and rather data rather than last years’ temperatures.

After decomposition, I performed **Single, Double and Triple (additive and multiplicative) Exponential smoothing**, and got the results, which can be found in **Step 5** of the solution (comparative table of 4 ES models).

Single ES did not consider trend and seasonality, and double ES considered trend, but did not take seasonality into account (which was obviously present, according to decomposition plot). In all cases, **alpha (base values) was closer to 1 rather than 0, suggesting that the model found it optimal to put more weight on recent observations rather than past ones and did not detect too much randomness in the system**. Baseline estimate ranged from 63 to 73 degrees. On the other hand, **beta (trend) values suggested that the trend slope of the data is not much dependent on recent trend slopes**. Final trend estimate was always almost 0, suggesting that **there is no significant trend in temperature change over 20 years**. Finally, **gamma values (seasonal component) for the triple ES suggested putting more weight on recent seasonalities**. For the triple ES models I also got seasonal factor estimates for each of the 123 days.

The final choice was between **Additive and Multiplicative Triple ES Models**. I fitted both models, plotted their decomposition and fit on the data. The difference was very little and rather hard to tell visually.

However, with **Additive seasonality** the **Sum of Squared Errors** was lower (66244 compared to 68904 with multiplicative seasonality), so **additive seasonality method was chosen**.

The logic for CUSUM was the following: use the known seasonal factors for each day of each year from the exponential smoothing model to build a CUSUM model for **each year** and detect violations. For each year, explore **which day** the seasonal factor **drops below the set margin** (lower violations). Then, **compare the dates of first violation** year-to-year to see if the dates are **getting later**. This would mean that the fall starts later each year, hence the summers are becoming longer. To do that, I put extracted from the model seasonal factors for each day of each year in a new data frame and used it for CUSUM model building.

To define the **No Change Period** for the model, based on which I would calculate means and standard deviations of seasonal factors for each year, I calculated **average seasonal factor value** for each day using all years' data and visualized it on a plot to see the general pattern of the factors over a season. I chose the first **51 days** as the no-change period (July 1 - Aug 20). A shorter period could be riskier in terms of false positives, as the factors for July only are in general high, and a longer period, on the other hand, could make the model miss real summer end dates due to variability in year-to-year seasonal factors (some extreme values could be included in the mean). 51 days period seemed like the perfect range of dates with a 'normal' factor values. With this period, the 19 (19 years, as 1996 is used as the first observation in exponential smoothing) calculated mean values of seasonal factors ranged from 4.5 to 5.8, and standard deviations from 3.5 to 4.6.

As for **C and T**, I decided to start with most common values (1 or 2 standard deviations for C, 4 or 5 standard deviations for T). I chose **C=1sd and T=4sd** (slightly lowered the violation margins to 4sd, making the model a bit more sensitive, since sd was not that little (3.5-4.6) compared to the range of average seasonal factor values (they go from 0 to 13 and from 0 to -17)). Since the last time unofficial summer end I found was around 20th September, I assumed that a good combination of C and T would produce a similar result ± 1 week.

The chosen C and T values turned out to be a **great combination**. As a result, I got unofficial fall start dates ranging from 20 to 24 September.

The answer to the question is negative, because there is no evidence that the summers have become longer - fall start dates vary year-to-year by one or two days, and there is no trend in such change. In fact, in 1997-1999 summer was a bit longer (fall starts on the 09.24) compared to 2014-2015 (09.22 and 09.21).

I went further and **compared the results with my previous CUSUM results**, where I calculated start of fall for each year **without exponential smoothing** using similar C and T values (C=1sd, T=5sd). The results have changed drastically - before exponential smoothing fall start dates range by almost 2 months, from August to October, which is why it was harder to answer the question whether the summers have gotten longer. Exponential smoothing helped to get accurate results without taking the excess 'white noise' and randomness into consideration.

Below is a step-by-step solution with more reasoning and explanation of each choice I made throughout this exercise.

Solution in R

Step 0: Load the libraries

```
library(dplyr)
library(tidyverse)
library(dslabs)
library(data.table)
library(ggplot2)
library(plotly)
```

```
library(outliers)
library(qcc)
```

Step 1: Load the dataset & do basic exploration

Here, I load the data set and rename the columns to remove the 'X' symbols near each year.

```
data <- read.table("temps.txt",
                  header = TRUE,
                  stringsAsFactors = FALSE,
                  sep = "",
                  dec = ".")

#rename to exclude X from att names
data <- data %>% rename_at(vars(starts_with("X")), funs(str_replace(., "X", "")))

## Warning: `funs()` was deprecated in dplyr 0.8.0.
## i Please use a list of either functions or lambdas:
##
## # Simple named list: list(mean = mean, median = median)
##
## # Auto named with `tibble::lst()`: tibble::lst(mean, median)
##
## # Using lambdas list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.


data[, "DAY"] <- as.Date(data[, "DAY"], "%d-%B")
#leave only month and day
data[, "DAY"] <- format(data[, "DAY"], "%m.%d")
head(data)

##      DAY 1996 1997 1998 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009
## 1 07.01   98   86   91   84   89   84   90   73   82   91   93   95   85   95
## 2 07.02   97   90   88   82   91   87   90   81   81   89   93   85   87   90
## 3 07.03   97   93   91   87   93   87   87   87   86   86   93   82   91   89
## 4 07.04   90   91   91   88   95   84   89   86   88   86   91   86   90   91
## 5 07.05   89   84   91   90   96   86   93   80   90   89   90   88   88   80
## 6 07.06   93   84   89   91   96   87   93   84   90   82   81   87   82   87
## 2010 2011 2012 2013 2014 2015
## 1    87    92   105   82   90   85
## 2    84    94    93   85   93   87
## 3    83    95    99   76   87   79
## 4    85    92    98   77   84   85
## 5    88    90   100   83   86   84
## 6    89    90    98   83   87   84
```

I will also perform some basic data exploration to learn more about the data.

```
boxplot(data[-1],
        xlab="Year",
        ylab="°F",
        col="lightblue",
        border="darkblue")
title(main="July-October Temperatures in Atlanta (1996-2015)")
```

July–October Temperatures in Atlanta (1996–2015)



There is no obvious trend of a changing median temperature over the years, according to the boxplot. It looks like temperatures in general were a bit higher in 2010 and 2011, but the median goes back down to previous values in 2013. Also, some years seem to be more ‘stable’ in temperature, and sometimes it varies more - there are years with short whiskers, as well as years with many outlying values. In general, I would suspect that the temperatures have not been increasing over the last few years. However, let’s do a thorough analysis with exponential smoothing to draw supported conclusions.

Step 2: Convert Data to Time Series

Since exponential smoothing is a time series forecasting method, we need to convert our data to time series.

The `ts()` function requires **two critical inputs - frequency and start**. Since we have years’ data from 1996 to 2015, **1996 would be our start**. And, because **we have 123 days’ temperature known for each year, 123 would be our frequency** - that is the number of rows in our data frame.

First, we need to merge all temperatures in a vector, and then generate a time-series object using the start and frequency values from above.

```
#There are two critical inputs we must give the function - frequency and start.

#Turn col data into one vector
data_vector <- as.vector(unlist(data[,2:ncol(data)]))

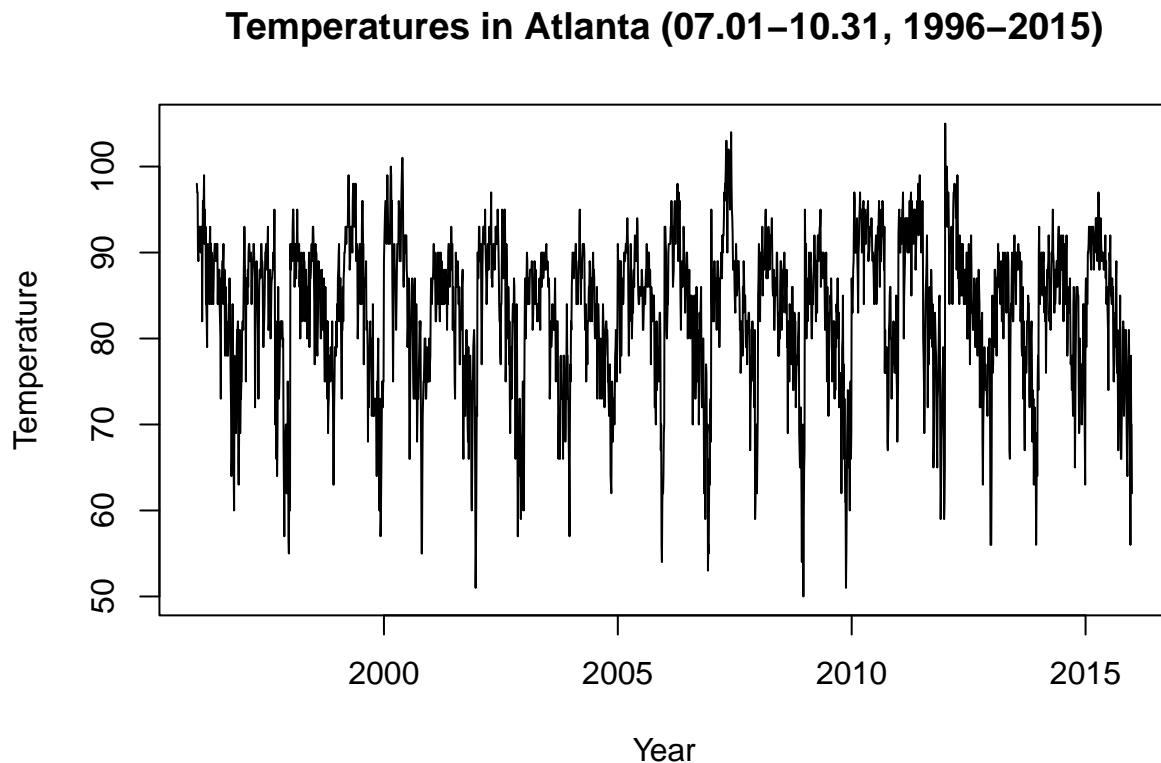
#Turn data to ts object
#we have 123 days' temp for each year - we will have 'seasons' from 1996 each lasting 123 days

datats <- ts(data_vector, start=1996, frequency=nrow(data))
head(datats)
```

```
## [1] 98 97 97 90 89 93
```

Visualize our time series object:

```
plot(datats, main="Temperatures in Atlanta (07.01-10.31, 1996-2015)", xlab="Year", ylab="Temperature")
```



Looking at the graph here, although there could be some factors influencing our data (for example, randomness), we can already see that there is no obvious trend - the 'arches' (our 'seasons') stay on the same level, and there is no distinctive upwards or downwards trend in the temperatures. Also, there are some 'anomalities' in each season - some very low or very high temperatures, but there is no obvious pattern connecting them.

Step 3: Decomposition

Let us try decomposition for our data to check how significantly factors like trend, seasonality and randomness affect our temperatures.


There are two types of decomposition we could perform - with **additive or multiplicative** seasonality.

Before trying both, let's think about what each seasonality type would mean for our data: -**Additive seasonality** would mean that our seasons remain similar in widths and heights, while there is a linear trend year-to-year. So, our average temperature would be increasing/decreasing year to year, but the amplitude of temperature variations would stay the same. -**Multiplicative seasonality**, on the other hand, would mean that our temperatures are becoming more extreme in their range year to year - for example, each year for the part of the season with warmer temperatures, the temperatures would become even more warmer, and for the part with lower temperatures of the season, unlike with additive seasonality, where they would become warmer as well (to keep the temperature spread similar to last year), they would become even colder.

Multiplicative seasonality would most probably make our data look like a funnel, so let's check the graph again to see if that is the case. I will also add a smoothing line to see the direction of temperature changes

```
ggplot(as.data.frame(datats), aes(x=c(1:2460), y=x)) +
  geom_line(aes(x=c(1:2460), y=x), color="darkblue") +
  geom_smooth(method="loess", color="green")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting to continuous.
## `geom_smooth()` using formula = 'y ~ x'
```




As we can see, our data does not look like a funnel. The seasonality is most probably additive. And, what is more important, it looks like there is no trend at all.

Keeping in mind that our seasonality factor would likely be additive, I will still decompose the data with both seasonality types to compare the results.

Decomposition (additive seasonality):

```
components_dts <- decompose(datats)
plot(components_dts)
```

Decomposition of additive time series



On the graph above besides the observed temperatures we see 3 other components of the data set: -**Trend**: shows the long-term trends in the data. **We can see that it barely ranges in values from 82 to 86, so as we suspected before, there is likely no significant long-term trend in the data**, hence we are unlikely to get proof that the summers are getting warmer or colder -**Seasonal**: repeated seasonal component that makes data vary with the seasons. **We have a significant seasonality factor - we can see how in the beginning of each season (summer) the temperatures are higher, and how they lower towards its end (October)** -**Random**: Components that are left over, and are not expected from seasonality or trends. **Our random component is quite significant - its range is bigger than the range of the seasonal factor or the trend**. Higher significance of randomness compared to seasonality and trend in general would make the prediction for the future period less accurate. That is logical - with almost no trend and this much randomness we would be unlikely to predict temperatures for each day in July-October 2016 using just past years' values. After all, there are lots of factors that influence weather, and it rather depends on the ongoing forces (like cyclones) than on the last years' data.

I will output the range of all factors to compare them to decomposition with multiplicative seasonality:

```
#trend
min(na.omit(components_dts$trend))

## [1] 80.45528

max(na.omit(components_dts$trend))

## [1] 88.02439

#seasonal
min(na.omit(components_dts$seasonal))

## [1] -15.23864
```

```
max(na.omit(components_dts$seasonal))
```

```
## [1] 7.728413
```

```
#random
```

```
min(na.omit(components_dts$random))
```

```
## [1] -21.59362
```

```
max(na.omit(components_dts$random))
```

```
## [1] 15.61819
```

Indeed, the random component has the widest scale, while the trend is insignificant (we are unlikely to find change in temperatures for the last years).

Decomposition (multiplicative seasonality):


Let us do the same with multiplicative seasonality:

```
#Decomposition (multiplicative)
```

```
components_dts_mult <- decompose(datats, type="mult")
```

```
plot(components_dts_mult)
```

Decomposition of multiplicative time series



The difference is that the random and seasonal components are now multiplicative instead of additive. In general, We still have a trend with low significance, good seasonality and lots of randomness.

Let's check the scale of each component:

```
#trend
```

```
min(na.omit(components_dts_mult$trend))
```



```
## [1] 80.45528
max(na.omit(components_dts_mult$trend))

## [1] 88.02439
#seasonal
min(na.omit(components_dts_mult$seasonal))

## [1] 0.8168494
max(na.omit(components_dts_mult$seasonal))

## [1] 1.092631
#random
min(na.omit(components_dts_mult$random))

## [1] 0.7041936
max(na.omit(components_dts_mult$random))

## [1] 1.209286
```

We have the same values for the trend component, and the Seasonal and Random factors are expressed differently (with coefficients).

Conclusion for step 3:

We now have found that it is unlikely for us to detect a trend in temperatures to answer the main question, whether the summers have become longer/warmer. It seems from the graphs that any change would more likely be due to change. Moreover, there is a seasonality component with good significance, so we would need to consider it in our final model with exponential smoothing. Finally, from visual inspection of the data it appear that we would rather have additive seasonality than multiplicative.

Now, let's move on to exponential smoothing.

Step 4: Exponential smoothing

Based on the findings in the previous step, we can say that the trend for our data is minimal, and temperature changes would probably be caused by randomness.

Now, let's apply exponential smoothing to our data. We will try different models using HoltWinters function and compare them.

We will then **extract the seasonal factors** for each day of the year and use them to perform CUSUM. Using these factors, we would be able to get rid of 'white noise' or randomness that could affect the accuracy of CUSUM.

There are **3 types of exponential smoothing** that we will try: -Single ES - simple ES that would use a weighted moving avg of our temperatures without trend and seasonality components -Double ES - a more reliable method that considers trends in the data, but no seasonality -Triple ES - a method for data with both trend and seasonality

From our previous findings, I would suspect that the model most suitable for our data would be **triple ES**, as we have both trend and seasonality components.

I will try all the models and compare the optimal values that they produce. I will also compare models with additive and multiplicative seasonality to make the final choice for the type of this component.

Single ES

```
#single - no trend no season
HW1 <- HoltWinters(datats,
                  beta=FALSE, #model without trend
                  gamma=FALSE) #not considering seasonality
```

```
HW1
```

```
## Holt-Winters exponential smoothing without trend and without seasonal component.
##
## Call:
## HoltWinters(x = datats, beta = FALSE, gamma = FALSE)
##
## Smoothing parameters:
##  alpha: 0.8388021
##  beta : FALSE
##  gamma: FALSE
##
## Coefficients:
##      [,1]
## a 63.30952
```

We have a baseline estimate of 63.309, and the optimal alpha as estimated by the model is **0.8388**. That means that there is not much randomness in the data, and we put more trust in the most recent observations (temperature in the most recent years).

Double ES

```
HW2 <- HoltWinters(datats,
                  gamma=FALSE) #now we have trend value in the M but still no season
```

```
HW2
```

```
## Holt-Winters exponential smoothing with trend and without seasonal component.
##
## Call:
## HoltWinters(x = datats, gamma = FALSE)
##
## Smoothing parameters:
##  alpha: 0.8445729
##  beta : 0.003720884
##  gamma: FALSE
##
## Coefficients:
##      [,1]
## a 63.2530022
## b -0.0729933
```

With an ES that includes trend as a component, we get a similar baseline estimate of 63, and a **slightly higher alpha of 0.8445** - so the amount of randomness is again found to be rather low, and the recent datapoints can be trusted.

The **trend value, beta** is close to zero, which means that **the trend slope of our temperatures is less dependent on the trend slopes for the recent temperatures**. In other words, recent temperature observations do not add to the trend value that much, and that resulted in a **final trend estimate close to 0, -0.07**. This one again is a signal that no significant trend is observed in our temperature observations over time, moreover, this small trend is negative, so the temperatures have not become warmer over the years.

Triple ES - Additive seasonality

First, let's try additive seasonality component gamma and check how it changes the results:

```
HW3_a <- HoltWinters(datats)
HW3_a

## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = datats)
##
## Smoothing parameters:
##   alpha: 0.6610618
##   beta : 0
##   gamma: 0.6248076
##
## Coefficients:
##              [,1]
## a      71.477236414
## b     -0.004362918
## s1     18.590169842
## s2     17.803098732
## s3     12.204442890
## s4     13.233948865
## s5     12.957258705
## s6     11.525341233
## s7     10.854441534
## s8     10.199632666
## s9       8.694767348
## s10     5.983076192
## s11     3.123493477
## s12     4.698228193
## s13     2.730023168
## s14     2.995935818
## s15     1.714600919
## s16     2.486701224
## s17     6.382595268
## s18     5.081837636
## s19     7.571432660
## s20     6.165047647
## s21     9.560458487
## s22     9.700133847
## s23     8.808383245
## s24     8.505505527
## s25     7.406809208
## s26     6.839204571
## s27     6.368261304
## s28     6.382080380
## s29     4.552058253
## s30     6.877476437
## s31     4.823330209
## s32     4.931885957
## s33     7.109879628
## s34     6.178469084
```

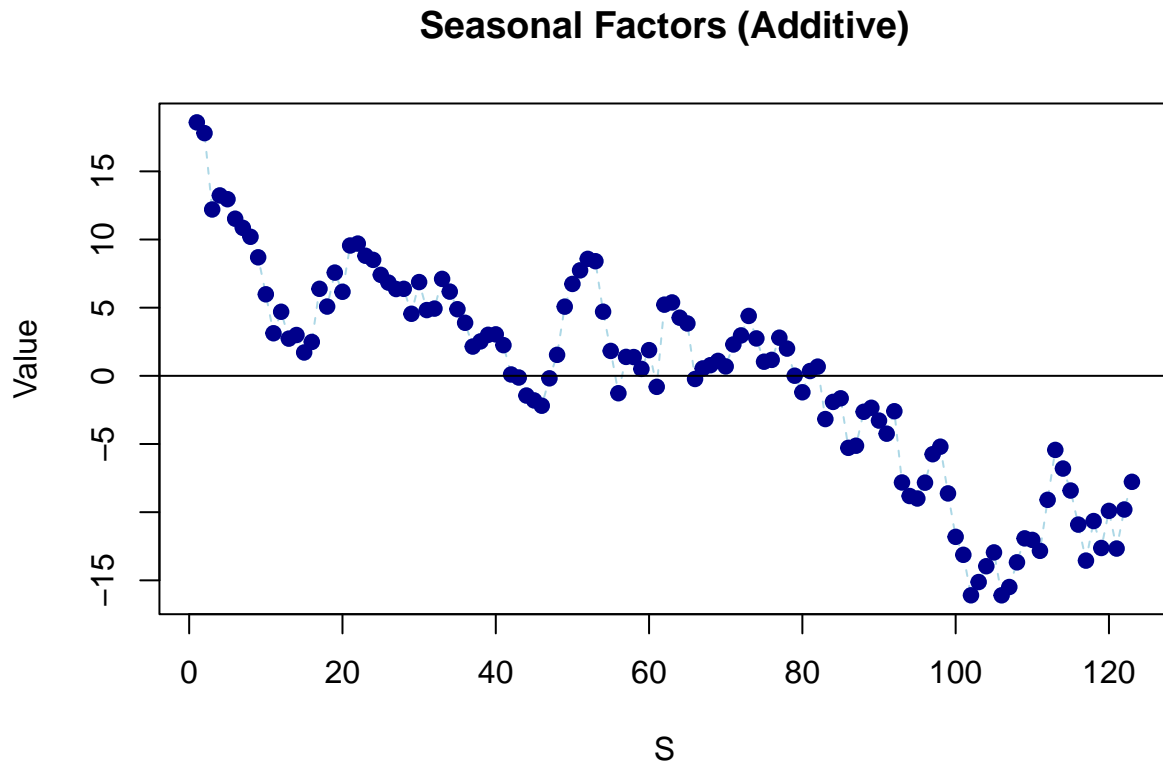
s35 4.886891317
s36 3.890547248
s37 2.148316257
s38 2.524866001
s39 3.008098232
s40 3.041663870
s41 2.251741386
s42 0.101091985
s43 -0.123337548
s44 -1.445675315
s45 -1.802768181
s46 -2.192036338
s47 -0.180954242
s48 1.538987281
s49 5.075394760
s50 6.740978049
s51 7.737089782
s52 8.579515859
s53 8.408834158
s54 4.704976718
s55 1.827215229
s56 -1.275747384
s57 1.389899699
s58 1.376842871
s59 0.509553410
s60 1.886439429
s61 -0.806454923
s62 5.221873550
s63 5.383073482
s64 4.265584552
s65 3.841481452
s66 -0.231239928
s67 0.542761270
s68 0.780131779
s69 1.096690727
s70 0.690525998
s71 2.301303414
s72 2.965913580
s73 4.393732595
s74 2.744547070
s75 1.035278911
s76 1.170709479
s77 2.796838283
s78 2.000312540
s79 0.007337449
s80 -1.203916069
s81 0.352397232
s82 0.675108103
s83 -3.169643942
s84 -1.913321175
s85 -1.647780450
s86 -5.281261301
s87 -5.126493027
s88 -2.637666754

```
## s89 -2.342133004
## s90 -3.281910970
## s91 -4.242033198
## s92 -2.596010530
## s93 -7.821281290
## s94 -8.814741200
## s95 -8.996689798
## s96 -7.835655534
## s97 -5.749139155
## s98 -5.196182693
## s99 -8.623793296
## s100 -11.809355220
## s101 -13.129428554
## s102 -16.095143067
## s103 -15.125436350
## s104 -13.963606549
## s105 -12.953304848
## s106 -16.097179844
## s107 -15.489223470
## s108 -13.680122300
## s109 -11.921434142
## s110 -12.035411347
## s111 -12.837047727
## s112 -9.095808127
## s113 -5.433029341
## s114 -6.800835107
## s115 -8.413639598
## s116 -10.912409484
## s117 -13.553826535
## s118 -10.652543677
## s119 -12.627298331
## s120 -9.906981556
## s121 -12.668519900
## s122 -9.805502547
## s123 -7.775306633
```

Once we added a seasonality component, gamma, our **alpha lowered to 0.66**, signaling that now the model considers to be more randomness than before, though the value is still closer to 1 rather than 0, so we rely more on the recent observations than on the past ones. This is not unexpected, since we saw that there was a random component when decomposing the data. The **baseline estimate is now 71** degrees, **beta is 0**, so the most recent trend slopes are not taken into account for the **final trend estimate of -0.004**. Once again, there is no trend in temperatures change in our data. Finally, we now have a **gamma, the seasonal component, of 0.62**, it is closer to 1 but not far from 0.5, so the weight of seasonal cycles in more or less even with a bit more accent on the recent seasonal cyclicities. We also have **123 seasonal factors**, containing final estimates of the seasonal component for each day from July 1st to October 31st.

```
plot(HW3_a$coefficients[3:length(HW3_a$coefficients)],
     type="o",
     col="lightblue",
     lty=2,
     lwd=1,
     main="Seasonal Factors (Additive)",
     xlab="S",
     ylab="Value")
points(HW3_a$coefficients[3:length(HW3_a$coefficients)],
```

```
col="darkblue",
pch=19, cex=1)
abline(h=0)
```



From the plot of the seasonal factor estimates we can see that around the 80th day of a season (~20th August) are negative-only. So after that period the temperatures are estimated to be below the baseline with an application of a corresponding seasonal factor.

Triple ES - Multiplicative seasonality

Let's check if we get a different result using multiplicative seasonality:

```
HW3_m <- HoltWinters(datats, seasonal = "multiplicative")
HW3_m
```

```
## Holt-Winters exponential smoothing with trend and multiplicative seasonal component.
##
## Call:
## HoltWinters(x = datats, seasonal = "multiplicative")
##
## Smoothing parameters:
##  alpha: 0.615003
##  beta : 0
##  gamma: 0.5495256
##
## Coefficients:
##              [,1]
## a      73.679517064
```

```
## b      -0.004362918
## s1      1.239022317
## s2      1.234344062
## s3      1.159509551
## s4      1.175247483
## s5      1.171344196
## s6      1.151038408
## s7      1.139383104
## s8      1.130484528
## s9      1.110487514
## s10     1.076242879
## s11     1.041044609
## s12     1.058139281
## s13     1.032496529
## s14     1.036257448
## s15     1.019348815
## s16     1.026754142
## s17     1.071170378
## s18     1.054819556
## s19     1.084397734
## s20     1.064605879
## s21     1.109827336
## s22     1.112670130
## s23     1.103970506
## s24     1.102771209
## s25     1.091264692
## s26     1.084518342
## s27     1.077914660
## s28     1.077696145
## s29     1.053788854
## s30     1.079454300
## s31     1.053481186
## s32     1.054023885
## s33     1.078221405
## s34     1.070145761
## s35     1.054891375
## s36     1.044587771
## s37     1.023285461
## s38     1.025836722
## s39     1.031075732
## s40     1.031419152
## s41     1.021827552
## s42     0.998177248
## s43     0.996049257
## s44     0.981570825
## s45     0.976510542
## s46     0.967977608
## s47     0.985788411
## s48     1.004748195
## s49     1.050965934
## s50     1.072515008
## s51     1.086532279
## s52     1.098357400
## s53     1.097158461
```

s54 1.054827180
s55 1.022866587
s56 0.987259326
s57 1.016923524
s58 1.016604903
s59 1.004320951
s60 1.019102781
s61 0.983848662
s62 1.055888360
s63 1.056122844
s64 1.043478958
s65 1.039475693
s66 0.991019224
s67 1.001437488
s68 1.002221759
s69 1.003949213
s70 0.999566344
s71 1.018636837
s72 1.026490773
s73 1.042507768
s74 1.022500795
s75 1.002503740
s76 1.004560984
s77 1.025536556
s78 1.015357769
s79 0.992176558
s80 0.979377825
s81 0.998058079
s82 1.002553395
s83 0.955429116
s84 0.970970220
s85 0.975543504
s86 0.931515830
s87 0.926764603
s88 0.958565273
s89 0.963250387
s90 0.951644060
s91 0.937362688
s92 0.954257999
s93 0.892485444
s94 0.879537700
s95 0.879946892
s96 0.890633648
s97 0.917134959
s98 0.925991769
s99 0.884247686
s100 0.846648167
s101 0.833696369
s102 0.800001437
s103 0.807934782
s104 0.819343668
s105 0.828571029
s106 0.795608740
s107 0.796609993



```
## s108 0.815503509
## s109 0.830111282
## s110 0.829086181
## s111 0.818367239
## s112 0.863958784
## s113 0.912057203
## s114 0.898308248
## s115 0.878723779
## s116 0.848971946
## s117 0.813891909
## s118 0.846821392
## s119 0.819121827
## s120 0.851036184
## s121 0.820416491
## s122 0.851581233
## s123 0.874038407
```

Now, we have a slightly **lower alpha and gamma, and the same beta=0**, so less weight is put on recent observations. **Baseline estimate is at 73.4**, and the trend estimate is again very insignificant and negative, **-0.004**.

Lets visualize seasonal factors for better understand of the difference of the two methods:

```
plot(HW3_m$coefficients[3:length(HW3_m$coefficients)],
     type="o",
     col="lightblue",
     lty=2,
     lwd=1,
     main="Seasonal Factors (Multiplicative)",
     xlab="S",
     ylab="Value")
points(HW3_m$coefficients[3:length(HW3_m$coefficients)],
       col="darkblue",
       pch=19, cex=1)
abline(h=1)
```

Seasonal Factors (Multiplicative)



Similar picture to the previous method - after 80th day index seasonality factors drag the temperature below the baseline.

Step 5: Choosing the model

Let's sum up the parameters of our models:

```
es_models <- data.frame(
  alpha=c(HW1$alpha,HW2$alpha,HW3_a$alpha,HW3_m$alpha),
  beta=c(NA,HW2$beta,HW3_a$beta,HW3_m$beta),
  gamma=c(NA,NA,HW3_a$gamma,HW3_m$gamma),
  Baseline_Estimate=c(HW1$coefficients[1],HW2$coefficients[1],HW3_a$coefficients[1],HW3_m$coefficients[1]),
  Trend_Estimate=c(HW1$coefficients[2],HW2$coefficients[2],HW3_a$coefficients[2],HW3_m$coefficients[2])

rownames(es_models) <- c('Single',
                        'Double',
                        'Triple (Additive)',
                        'Triple (Multiplicative)')

es_models
```

	alpha	beta	gamma	Baseline_Estimate
## Single	0.8388021	NA	NA	63.30952
## Double	0.8445729	0.003720884	NA	63.25300
## Triple (Additive)	0.6610618	0.000000000	0.6248076	71.47724
## Triple (Multiplicative)	0.6150030	0.000000000	0.5495256	73.67952
##	Trend_Estimate			

```
## Single NA
## Double -0.072993299
## Triple (Additive) -0.004362918
## Triple (Multiplicative) -0.004362918
```


As we can see, once we introduce seasonality to the model, the alpha becomes lower and not as much weight is put on the most recent observations, as with the previous two models.

The trend estimate is close to zero and negative for all models where it is taken into account. **This signals that there is no change in our temperatures throughout the years**, so the answer to the main question is likely to be negative. Finally, multiplicative model put a bit less weight on most recent seasonal cycles and provides a slightly higher baseline estimate.

We should be choosing between additive and multiplicative triple ES models, as they take into account both trend and seasonality, and the latter is very significant for our temperatures data.


First, let's compare the decomposition plots of both models:

```
plot(fitted(HW3_a), main="Decomposition (Additive)")
```



```
plot(fitted(HW3_m), main="Decomposition (Multiplicative)")
```

Decomposition (Multiplicative)



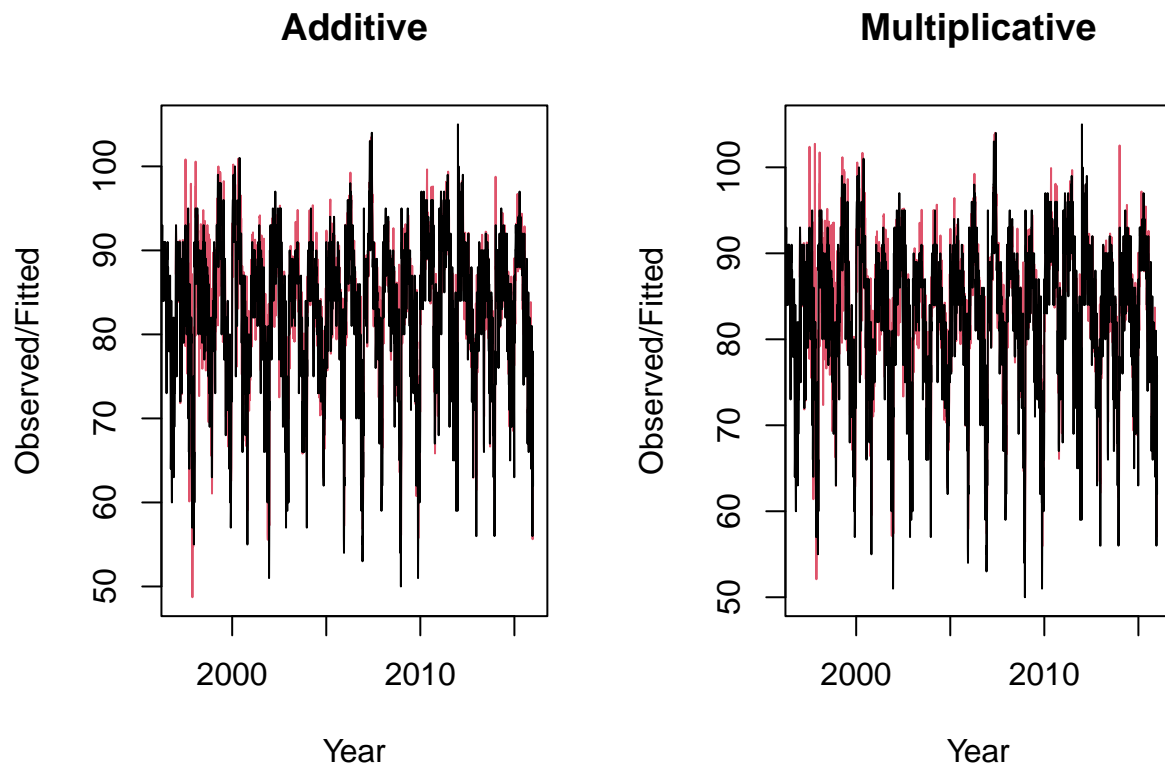
The results are very similar - both models estimate a trend close to zero, and seasonalities have year-to-year repetitive patterns, except for the last few years in the multiplicative model, where the range of the seasonal factor becomes slightly wider (due to high seasonality factors at the beginning of the season).

This is not enough to make the final choice, so we keep exploring both models.

Let's try and visualize how both models fit the data:

```
par(mfrow=c(1,2))
plot(HW3_a, main="Additive", xlab="Year", ylab="Observed/Fitted")

plot(HW3_m, main="Multiplicative", xlab="Year", ylab="Observed/Fitted")
```



The first year is not present for the fitted values, as they are used as a first observation in ES and are not fitted according to the ES equation. Both models have a very similar fit, the only difference that is noticeable visually are the slightly higher fitted values for the first few years in the multiplicative model.

Finally, let's compare the Sum of Squared Errors for each model

```
print('SSE Additive')
```

```
## [1] "SSE Additive"
```

```
HW3_a$SSE
```

```
## [1] 66244.25
```

```
print('SSE Multiplicative')
```

```
## [1] "SSE Multiplicative"
```

```
HW3_m$SSE
```

```
## [1] 68904.57
```

SSE is smaller for the Additive model. Hence, we should move on with the Additive ES - our initial assumption (based on visual inspection) that additive seasonality suits the data better was correct.

Step 6: Preparation for CUSUM

Although we have found the trend estimate to be around 0, we should still check if there is any change in the seasonal factor values for the past few years. On the decomposition plot we have seen that the seasonality factor does not significantly change for the latest years - the length of seasons does not increase (the 'arches'

on the graph are not getting wider year to year), and the duration of warm parts of each season (the width of the higher part of the arch) appear to be the same for each year. However, it is worth checking with CUSUM for any possible change in seasonal factors, as they may be hard to detect using visual inspection.

The logic for our CUSUM model will be the following: Since we know seasonal factors for each day of each year from our ES model with additive seasonality, we can use them to detect whether the unofficial end of summer has become later. We will run CUSUM on seasonal factors for each data point for each year. By doing this we can see where in each year (on which day) a change was detected by CUSUM, and then we will compare those dates to see if the change happens later over time.

Let's have a look at our fitted model output:

```
head(HW3_a$fitted)

##           xhat      level      trend      season
## [1,] 87.17619 82.87739 -0.004362918  4.303159
## [2,] 90.32925 82.09550 -0.004362918  8.238119
## [3,] 92.96089 81.87348 -0.004362918 11.091777
## [4,] 90.93360 81.89497 -0.004362918  9.042997
## [5,] 83.99752 81.93450 -0.004362918  2.067387
## [6,] 84.04358 81.93177 -0.004362918  2.116168
```

We need the 'season' column for the CUSUM - it contains the needed season factors for each day in the data set.

Now let's transform those factors to a matrix and rename columns and rows so that we can use it for CUSUM analysis.

```
#put seasonal factors in the matrix
season_factors <- matrix(HW3_a$fitted[,4],nrow=123)
dim(season_factors)
```

```
## [1] 123  19
```

```
head(season_factors)

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,]  4.303159  4.054077  9.349191  8.130519  9.261932  8.532675 10.833750
## [2,]  8.238119  8.168393  8.457426  7.810603  8.686298  9.197265  9.837370
## [3,] 11.091777 11.100059 11.213419 11.470326 11.416575 11.012491 10.210645
## [4,]  9.042997  9.057058  9.529053 10.185524 10.863856 10.209557 10.532285
## [5,]  2.067387  2.067912  3.708913  5.588420  7.004562  8.024549  9.444664
## [6,]  2.116168  2.106939  2.232254  3.394700  4.340175  5.462823  6.487578
##           [,8]      [,9]      [,10]     [,11]     [,12]     [,13]     [,14]
## [1,] 10.596553  9.559008  9.526730 12.454249 15.18901 15.052357 17.963213
## [2,] 11.663073 10.874704 10.162519 11.021049 10.13469 11.583196 12.246533
## [3,] 12.021948 12.738520 11.468041 11.483482 10.45074 11.722773 11.707197
## [4,] 10.867244 11.779117 11.552601 11.117308 11.69283 11.650082 12.084663
## [5,]  8.518818  9.749682 10.738920 10.552850 11.29192 10.939712  8.908885
## [6,]  7.648008  8.250531  7.421824  6.156207  7.12694  6.619883  8.329696
##           [,15]     [,16]     [,17]     [,18]     [,19]
## [1,] 17.946705 17.796582 22.385105 21.31072 19.45438
## [2,] 12.817176 14.277046 14.037278 16.07719 17.19256
## [3,] 11.803979 12.726046 14.244787 12.98723 12.74993
## [4,] 12.461996 12.000781 12.458283 12.62307 11.98537
## [5,] 10.345553 10.214813 11.172561 12.77225 12.94898
## [6,]  9.151988  9.361362  9.444097 10.35325 11.13812
```

We have 19 year (1997 to 2015, as 1996 is the first observation and is not included) and 123 days in each.

Renaming columns and rows to make data look as the initial data set for convenience:

```
#rename
colnames(season_factors) <- c(1997:2015)
rownames(season_factors) <- data$DAY
#check
head(season_factors)
```

```
##           1997      1998      1999      2000      2001      2002      2003
## 07.01  4.303159  4.054077  9.349191  8.130519  9.261932  8.532675 10.833750
## 07.02  8.238119  8.168393  8.457426  7.810603  8.686298  9.197265  9.837370
## 07.03 11.091777 11.100059 11.213419 11.470326 11.416575 11.012491 10.210645
## 07.04  9.042997  9.057058  9.529053 10.185524 10.863856 10.209557 10.532285
## 07.05  2.067387  2.067912  3.708913  5.588420  7.004562  8.024549  9.444664
## 07.06  2.116168  2.106939  2.232254  3.394700  4.340175  5.462823  6.487578
##           2004      2005      2006      2007      2008      2009      2010
## 07.01 10.596553  9.559008  9.526730 12.454249 15.18901 15.052357 17.963213
## 07.02 11.663073 10.874704 10.162519 11.021049 10.13469 11.583196 12.246533
## 07.03 12.021948 12.738520 11.468041 11.483482 10.45074 11.722773 11.707197
## 07.04 10.867244 11.779117 11.552601 11.117308 11.69283 11.650082 12.084663
## 07.05  8.518818  9.749682 10.738920 10.552850 11.29192 10.939712  8.908885
## 07.06  7.648008  8.250531  7.421824  6.156207  7.12694  6.619883  8.329696
##           2011      2012      2013      2014      2015
## 07.01 17.946705 17.796582 22.385105 21.31072 19.45438
## 07.02 12.817176 14.277046 14.037278 16.07719 17.19256
## 07.03 11.803979 12.726046 14.244787 12.98723 12.74993
## 07.04 12.461996 12.000781 12.458283 12.62307 11.98537
## 07.05 10.345553 10.214813 11.172561 12.77225 12.94898
## 07.06  9.151988  9.361362  9.444097 10.35325 11.13812
```

```
tail(season_factors)
```

```
##           1997      1998      1999      2000      2001      2002
## 10.26 -6.4935885 -6.890605 -5.6558405 -4.4347630 -5.3063651 -7.98863484
## 10.27 -6.4854584 -9.585617 -8.5954582 -6.2875108 -5.5543340 -9.16304674
## 10.28 -0.5342389 -3.267917 -4.2692969 -4.8258201 -4.4623961 -5.06875912
## 10.29  0.4332408  1.208673 -0.2895929 -1.5318538 -3.1641748 -1.94952674
## 10.30  0.4088506  1.101306  1.4642322  1.3080103  0.5778192  2.10392807
## 10.31 -0.6480600 -1.459240 -1.4283706 -0.8678262 -0.4418413  0.08050285
##           2003      2004      2005      2006      2007      2008
## 10.26 -6.7774517 -8.4837323 -8.1417726 -7.33985781 -6.7622484 -6.67570860
## 10.27 -8.2911287 -9.3950601 -8.6616976 -7.85532651 -6.7023824 -6.89657593
## 10.28 -4.3687338 -7.0550231 -7.9364006 -8.02753487 -8.2346789 -7.33976472
## 10.29 -2.3719010 -0.9513094 -3.1770223 -3.79134413 -2.8517916 -5.38165712
## 10.30 -0.3791102  1.1636893  0.5976518  0.64984731  0.2404763 -0.21206523
## 10.31 -2.0258362 -1.5768118 -0.3402956  0.08872082 -0.9891323 -0.03411151
##           2009      2010      2011      2012      2013      2014
## 10.26 -6.7927592 -8.4238940 -7.9411253 -8.7304960 -9.50699113 -9.707349
## 10.27 -9.2180348 -10.3152173 -11.6560776 -10.7124392 -12.67269269 -10.375104
## 10.28 -9.9377581 -7.8268114 -9.2308714 -10.4825572 -14.15951727 -14.123842
## 10.29 -4.7700225 -4.3009108 -6.7931879 -10.9093072 -12.06427234 -10.165486
## 10.30  0.1720216 -2.6205780 -3.1849145 -4.9196868 -6.57865384 -6.460541
## 10.31  0.4828511  0.1067451  0.1859101 -0.2679054  0.09155862 -2.127128
##           2015
## 10.26 -8.684994
## 10.27 -9.886253
```

```
## 10.28 -14.645752
## 10.29 -12.026808
## 10.30 -8.933947
## 10.31 -4.941084
```

Everything looks correct, now we can do our CUSUM Analysis.

Step 7: CUSUM

As usual, before we perform CUSUM, we need to define our **no change period**.

Last time I used a period from 1st to 31st July.

Let's plot average seasonal factor values for each day to check if we can use the same period:

```
sf_avg<-data.frame(rowMeans(season_factors, n = 19))
colnames(sf_avg)<- 'Avg_seasonal_factor'
head(sf_avg)
```

```
##      Avg_seasonal_factor
## 07.01      12.826311
## 07.02      11.183289
## 07.03      11.769472
## 07.04      11.141773
## 07.05       8.740071
## 07.06       6.691713
```

```
#ggplot(sf_avg, aes(x=1:123, y=Avg_seasonal_factor))+
#  geom_line()+
#  geom_point()
```

```
plot_ly(data=sf_avg, x=1:123, y=~Avg_seasonal_factor, type='scatter', mode='lines+markers', fill='tozero')
```

Based on the graph, we can see that until day 40 (Aug 9), there were no negative values for seasonal factors. On day 44 factor values go back up, but are not as high as before.

Since the values would vary for each year, I would try to choose a safe range of first 51 days (1 July - 20 August). If we choose values up to day 60, for example, we could get more negative seasonal factor values, which might lower our mean but also increase standard deviation for some years. By increasing the period we risk adding 'summer end' values, if there are years when summer ended in the last week of August. Period of 51 days includes enough variability of the factors (they are not as stable in growth over the 123 days as temperatures were the last time) to calculate the optimal 'normal' values for each years' model. Moreover, from the last HW3 I remember that summer never ended before the last week of August (and more often in fall), so we are not including any values for fall here.

Moving on to the CUSUM, let's create vectors to store CUSUM models and violations for each year.

```
models_cusum <- vector(mode="list", length=ncol(season_factors))
violat_cusum <- vector(mode="list", length=ncol(season_factors))
```

Since we chose July as the no change period, we need to find the mean and standard deviation values for each year's seasonal factors:

```
#mean and stdev of seasonal factors for each year
#use July temps
mean_seasons <- vector(mode="list", length=0)
stdev_seasons <- vector(mode="list", length=0)
for (i in 1:ncol(season_factors)){
  mean_seasons <- append(mean_seasons,mean(season_factors[1:51,i]))
```



```

stdev_seasons <- append(stdev_seasons, sd(season_factors[1:51, i]))
}
unlist(mean_seasons)

## [1] 4.561409 4.559721 4.651009 4.692767 4.808380 4.833226 4.939809 5.058915
## [9] 5.095669 5.114876 5.205778 5.321256 5.392601 5.482570 5.582841 5.654167
## [17] 5.727204 5.796219 5.816636

unlist(stdev_seasons)

## [1] 4.543606 4.534205 4.290408 4.145315 4.076333 3.901140 3.916727 3.839413
## [9] 3.796743 3.502302 3.587251 3.619594 3.711506 3.802659 3.898234 3.990965
## [17] 4.367116 4.562904 4.477043

```

We have mean values ranging from 4.5 to 5.8 and st.dev. from 3.5 to 4.6. These values do not change so much year to year, which is good for our model and suggests that we chose the no change period right.

I will be using the `cusum()` function for each year. To start with, I will set `C` and `T` to standard values. `C` is usually set to 1 or 2 standard deviation, and `T` to 4 or 5 st.dev. I will set `C` to 1sd and `T` to 4 sd (since our st. devs are between 3.5 and 4.6, which is not that small considering the scale of our season factors, which go from 0 to 13 and 0 to -17). Then based on the results I will adjust those values to balance the sensitivity of the model.

```

dec_int <- 4 #4 st deviations for T
se_shift <- 2 #2 instead of 1, since shift = 2C. SO 1st deviation for C

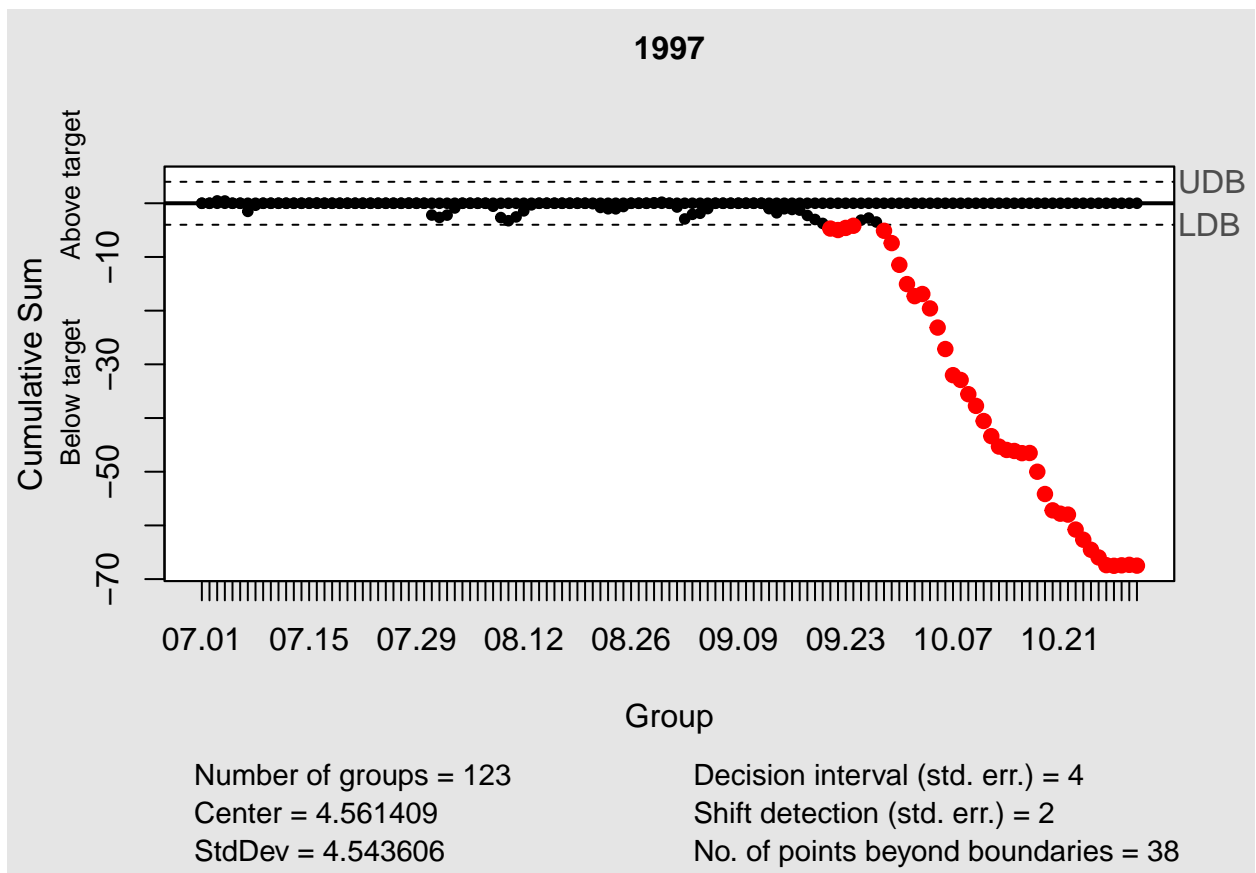
```


Finally, let's run CUSUM and see if we detect any violations


```

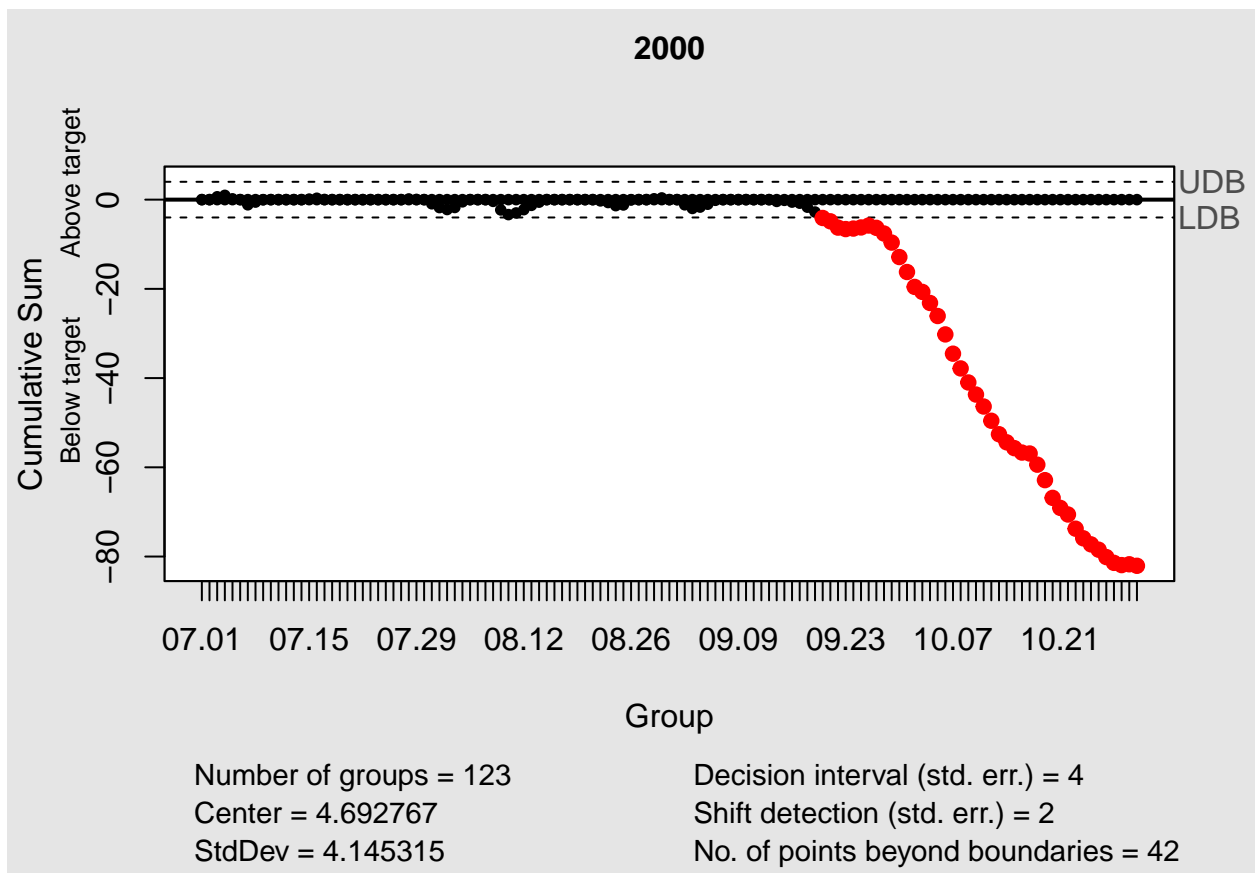
for (y in 1:ncol(season_factors)) {
  #cusum model for each year
  models_cusum[[y]] <- cusum(season_factors[,y], #column of the year in loop
                             center = unlist(mean_seasons)[y],
                             std.dev = unlist(stdev_seasons)[y], #use st dev for the particular year
                             decision.interval = dec_int, #T value, 5 st dev-s
                             se.shift = se_shift, #shift value, 2C, 2
                             plot = TRUE,
                             title = colnames(season_factors)[y]) #add plot for each year's cusum
  #cusum violation for each year - moment when fall came
  #use the model we built above
  #$use lower violation - the first day that violated temp - the 1st day of fall!
  violat_cusum[[y]] <- models_cusum[[y]]$violations$lower}


```

























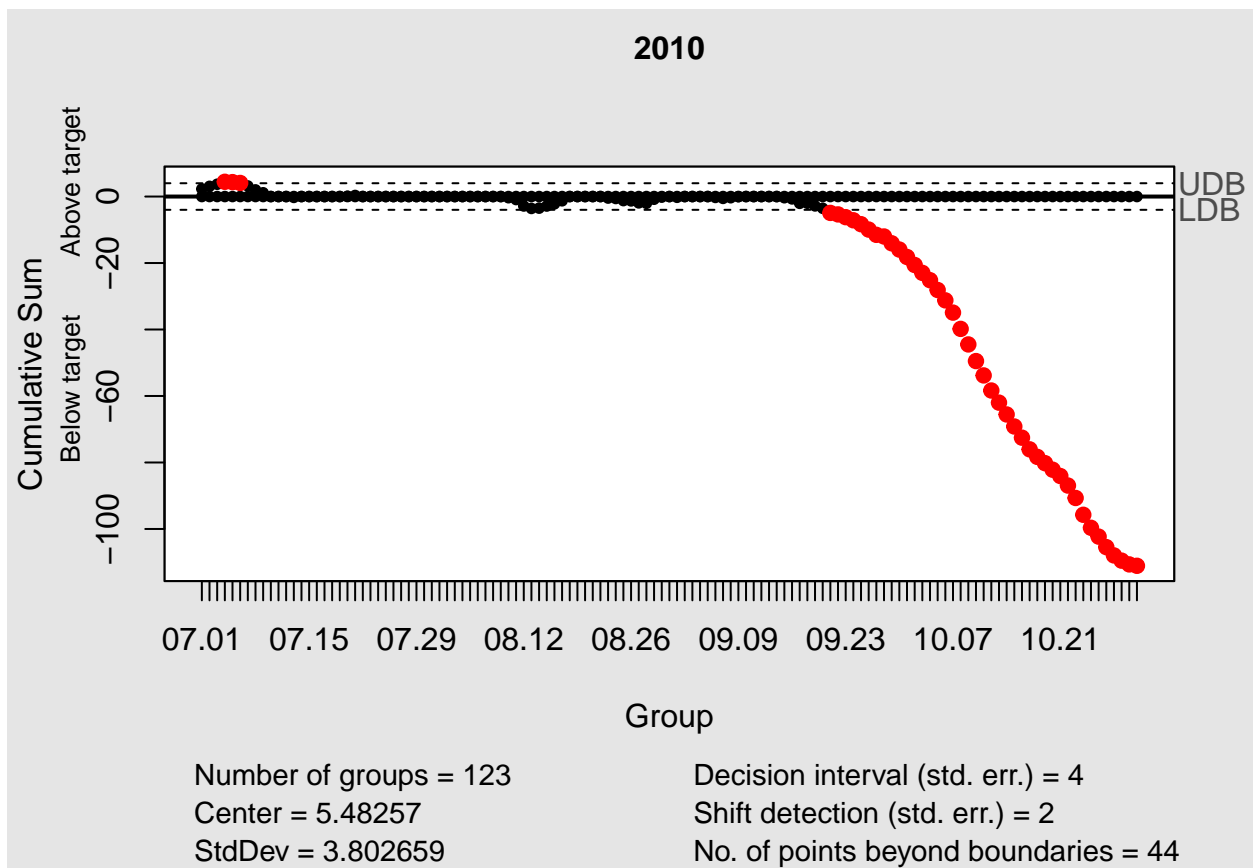



























It seems like the choice for C and T was good. **We are looking for the lower violations on the graphs** - it is obvious that the change is usually detected after the 3rd week of September - a similar result compared to the one I got in the CUSUM project.

Visually, it **does not appear that the summer is ending later**. But let's see the exact dates for each year:

```
#vector to store first days of fall
first_days <- vector(mode="list", length=ncol(season_factors))

#add to that vector first days
for (y in 1:ncol(season_factors)){
  first_days[y] <- min(violat_cusum[[y]])
}

#day index of first day of fall
first_days <- unlist(first_days)

#turn indexes to days
fall_dates<- vector(mode="list", length=0)
for (x in 1:19){
  i=0
  d<-first_days[x]
  fall_dates <- append(fall_dates, data[d,1], after=i)
  i=i+1
}

fall_dates<-unlist(fall_dates)
```

```
fall_dates
```

```
## [1] "09.24" "09.24" "09.23" "09.22" "09.22" "09.21" "09.21" "09.21" "09.20"  
## [10] "09.20" "09.21" "09.21" "09.21" "09.20" "09.20" "09.20" "09.21" "09.22"  
## [19] "09.21"
```

We can see that the date for first violation does not vary significantly, and there is no increasing trend. To see that clearly, let's visualize the results.

Data frame of unofficial fall start dates:


```
year_list <- 1997:2015  
year_list <- as.list(year_list)  
date_year <- do.call(rbind, Map(data.frame, Year=year_list, FallStart=fall_dates))  
date_year
```

```
##      Year FallStart  
## 1  1997      09.24  
## 2  1998      09.24  
## 3  1999      09.23  
## 4  2000      09.22  
## 5  2001      09.22  
## 6  2002      09.21  
## 7  2003      09.21  
## 8  2004      09.21  
## 9  2005      09.20  
## 10 2006      09.20  
## 11 2007      09.21  
## 12 2008      09.21  
## 13 2009      09.21  
## 14 2010      09.20  
## 15 2011      09.20  
## 16 2012      09.20  
## 17 2013      09.21  
## 18 2014      09.22  
## 19 2015      09.21
```

Plot:

```
results_plot<- ggplot(data = date_year, aes(x = Year, y = FallStart, group=1)) +  
  geom_line(linetype="dashed", color="darkblue") +  
  geom_hline(yintercept = 2, color="lightblue")+ #compare to 09.21 for 2015  
  geom_point(color="blue")+  
  labs(title="Fall Start Dates in Atlanta (1997-2015)", y="Date")  
results_plot
```

Fall Start Dates in Atlanta (1997–2015)



As we can see, **there is no evidence that the summers have become longer**. Compared to 1997-1999, it even seems that summers have become slightly shorter. Hence, **the answer to our main question is No**.

Step 8: Extra - comparing results with and without smoothing

I would like to compare the results to my CUSUM project where I defined fall start for each year without exponential smoothing. I will not repeat CUSUM here, and just copy the start dates for each year. These are result for $C=1sd$ and $T=5sd$.


```
fall_dates_hw3 <- c('09.25', '08.16', '08.15', '09.05', '09.28', '09.01', '08.14', '09.18', '09.09', '10.07', '08.16')
year_list1 <- 1997:2015
year_list1 <- as.list(year_list1)
date_year1 <- do.call(rbind, Map(data.frame, Year=year_list, FallStart=fall_dates_hw3))
date_year1
```

```
##      Year FallStart
## 1  1997      09.25
## 2  1998      08.16
## 3  1999      08.15
## 4  2000      09.05
## 5  2001      09.28
## 6  2002      09.01
## 7  2003      08.14
## 8  2004      09.18
## 9  2005      09.09
## 10 2006     10.07
```

```
## 11 2007      08.10
## 12 2008      09.10
## 13 2009      08.31
## 14 2010      09.02
## 15 2011      09.06
## 16 2012      09.22
## 17 2013      08.15
## 18 2014      09.25
## 19 2015      09.02
```

Plot previous results and compare with new ones:

```
previous<-ggplot(data = date_year1, aes(x = Year, y = FallStart, group=1)) +
  geom_line(linetype="dashed", color="darkgreen") +
  geom_point(color="green")+
  labs(title="No Exp.Sm. - Fall Start Dates", y="Date")
previous
```



The effect of exponential smoothing is indeed drastic. Before it, fall start dates varied from Aug 10 to Oct 07 - almost 2 months' difference compared to just 4 days (09.20-09.24) with exponential smoothing.

To conclude, Exponential smoothing has helped us get rid of the excess 'noise' and randomness in data, showed that there is no trend in temperature changes year to year, and proved that the summers are not getting longer.